

INTEL UNNATI INDUSTRIAL TRAINING PROGRAM-2024
PROBLEM STATEMENT-4

***Introduction to GenAI and
Simple LLM inference on
CPU and finetuning of LLM
Model to create a Custom
Chatbot***

BY

TEAM BINARY BOLTZ (MRCET)

THUMMALAPALLI CHANDINI – 21N31A66H9

SATHURI TANMAI – 21N31A66F6

YETUKURI KARTHIKAY – 21N31A66K4

THIRUNAHARI VARSHINI – 21N31A66H7

KOTIPALLI ADITHI NAIDU – 21N31A6692

Under the Guidance of:

Dr. A.V.H. SAI PRASAD

Assoc. Professor

Department of Computational Intelligence

Head of the Department

Dr. D. SUJATHA



PROJECT INFORMATION

Problem Statement No.	PS-4
Problem statement	Introduction to GenAI and Simple LLM inference on CPU and finetuning of LLM Model to create a Custom Chatbot

BINARY BOLTZ TEAM

Team Leader	Thummalapalli Chandini
Team Members	Sathuri Tanmai Yetukuri Karthikay Thirunahari Varshini Kotipalli Adithi Naidu
Institution	Malla Reddy College of Engineering and Technology

MENTORS

Faculty Mentor	Dr. A.V.H. Sai Prasad
Industry Mentor	Ms. Vasudha Kumari
External Mentor	Mr. Abhishek Nandy

INTEL UNNATI INDUSTRIAL TRAINING PROGRAM – 2024

CERTIFICATE

This is to certify that the project submitted by the team Binary Boltz has been approved and successfully completed as part of the Intel Unnati Industrial Training Program - 2024. The project has been built as a response to PS-4: Introduction to GenAI - Simple LLM Inference on CPU and Fine-Tuning of LLM Model to Create a Custom Chatbot.

Team Name: Binary Boltz

Team Leader Name: Thummalapalli Chandini

Team Members:

- i Thummalapalli Chandini
- ii Sathuri Tanmai
- iii Yetukuri Karthikay
- iv Thirunahari Varshini
- v Kotipalli Aditi Naidu

Dr. A.V.H. SAI PRASAD

Associate Professor

Department of Computational Intelligence

Malla Reddy College of Engineering and Technology

INDEX

1.ABSTRACT	3
2.INTRODUCTION	4
3.SYSTEM REQUIREMENTS	6
4.IMPLEMENTATION	7
5.CONCLUSION	16

1.ABSTRACT

This report chronicles the development and optimization journey of a Large Language Model (LLM) focused on generating imaginative text. The primary goal is to empower the model, named Dolly-v2-3b, to produce compelling and coherent narratives in response to diverse prompts. Through fine-tuning on a specialized dataset consisting of 15,000 instruction/response pairs from various domains, the model excels in tasks such as creative text generation. Its efficiency, supported by 3 billion parameters, ensures high-quality outputs while maintaining computational efficiency, which is crucial for practical applications prioritizing responsiveness and cost-effectiveness.

Integration with the Intel Extension for Transformers significantly enhances the model's performance. This collaboration optimizes hardware usage, resulting in quicker inference times and improved efficiency during text generation tasks. Evaluation metrics such as eval_loss and eval_ppl underscore the model's accuracy and predictive capability, highlighting its ability to deliver precise and contextually appropriate responses.

Benchmarking exercises underscore the model's robustness, with metrics indicating minimal latency and high throughput during inference. For instance, the model processes 100 samples in approximately 14.16 seconds, achieving an average throughput of 7.061 samples per second. This showcases its suitability for real-time applications that demand rapid response capabilities. Additionally, this report explores the impact of fine-tuning methodologies, employing a structured approach to ensure the model's outputs adhere to ethical standards and inclusivity. By incorporating prompts that encourage socially conscious storytelling, the training process mitigates bias and fosters the creation of engaging, unbiased narratives.

Keywords: Large Language Models, Fine-Tuning, Alpaca Dataset, CPU Inference.

2.INTRODUCTION

2.1. Objective:

The objective of this report is to provide a comprehensive introduction to Generative AI (GenAI) and the processes involved in performing simple inference using Large Language Models (LLMs) on CPU resources. Additionally, this report aims to detail the methodologies for fine-tuning LLMs to develop custom chatbots tailored to specific applications. Through this, we will explore the foundational concepts of GenAI, demonstrate the practical steps for running LLM inferences on CPU, and outline the procedures for fine-tuning LLMs, ultimately enabling the creation of specialized, high-performance chatbot solutions.

2.2. Problem Statement:

Introduction to GenAI and Simple LLM inference on CPU and finetuning of LLM Model to create a Custom Chatbot.

2.3. Background of project:

The project aims to leverage the advancements in Generative AI (GenAI) and Large Language Models (LLMs) to develop a custom chatbot using the NousResearch Llama-2-7b-chat-hf model. GenAI and LLMs have revolutionized natural language processing by enabling models to understand and generate human-like text with high accuracy. This has significant applications in various industries, including customer service, healthcare, and education.

This project focuses on optimizing the performance of LLMs for environments with limited computational resources by performing inference on CPUs and using techniques like quantization and fine-tuning. Fine-tuning involves adapting a pre-trained model to a specific domain or dataset, improving its performance on specialized tasks.

The chosen dataset for fine-tuning is "mlabonne/guanaco-llama2-1k," which provides domain-specific conversational data. The project employs Quantized Low-Rank Adaptation (QLoRA) parameters to enhance the model's efficiency and performance. By fine-tuning the Llama-2-7b-chat-hf model and applying these optimizations, the project aims to create an effective, high-performance chatbot tailored to specific applications.

Ultimately, this project seeks to demonstrate the practical application of GenAI and LLMs in developing intelligent, resource-efficient chatbot solutions that can address specific needs and enhance user interactions across various domains.

2.4. Scope of the project:

The scope of this project involves developing and fine-tuning a custom chatbot using the NousResearch Llama-2-7b-chat-hf model. It includes selecting and preparing a suitable dataset for domain-specific adaptation and optimizing the model for efficiency. The project covers configuring training parameters, conducting the fine-tuning process, and evaluating the fine-tuned model. Additionally, it involves creating a standalone model ready for deployment to provide intelligent, domain-specific conversational capabilities in various applications.

2.5. Project features:

- **Advanced Model Selection:**

Utilizes the NousResearch Llama-2-7b-chat-hf, a state-of-the-art Large Language Model (LLM) known for its robust natural language understanding and generation capabilities.

- **Domain-Specific Fine-Tuning:**

Fine-tunes the base model on the "mlabonne/guanaco-llama2-1k" dataset to adapt it for specific conversational domains, enhancing its relevance and performance in targeted applications.

- **Model Optimization:**

Implements Quantized Low-Rank Adaptation (QLoRA) techniques, including 4-bit quantization, to optimize the model for efficiency, making it suitable for environments with limited computational resources.

- **Performance Evaluation:**

Evaluates the fine-tuned model using sample prompts to ensure it meets the desired performance standards and provides accurate, context-aware responses.

- **Standalone Model Creation:**

Merges the fine-tuned model with the base model to create a standalone version, ready for deployment without dependency on external components.

3. SYSTEM REQUIREMENTS

3.1. Hardware Requirements:

- Cloud Services
- RAM
- T4 GPU
- Network Connection

3.2. Software Requirements:

- **Operating System:**
Linux (Ubuntu 18.04 or later) or Windows 10/11.
- **Environment:**
Google Colab
- **Python:**
Python 3.8 or later.
- **Required Libraries and Frameworks:**
PyTorch
Transformers
PEFT
BitsAndBytes
TRL
- **Additional Tools:**
TensorBoard

4. IMPLEMENTATION

4.1. Source code:

```
"""
installing required libraries

!pip install pyarrow==14.0.1
!pip install requests==2.31.0 (for version incompatibility issue)

!pip install -q accelerate==0.21.0 peft==0.4.0 bitsandbytes==0.40.2 transformers==4.31.0
trl==0.4.7

"""
#####

#installing Required Libraries

import os
import torch
from datasets import load_dataset
from transformers import (
    AutoModelForCausalLM,
    AutoTokenizer,
    BitsAndBytesConfig,
    HfArgumentParser,
    TrainingArguments,
    pipeline,
    logging,
)
from peft import LoraConfig, PeftModel
from trl import SFTTrainer

#####

#assigning parameters

#Base Model
model_name = "NousResearch/Llama-2-7b-chat-hf"

# DataSet for FineTuning
dataset_name = "mlabonne/guanaco-llama2-1k"

# Fine-tuned model name
new_model = "Llama-2-7b-chat-finetune"
```

```
# QLoRA parameters
# LoRA dimension
lora_r = 64

# Alpha parameter
lora_alpha = 16
lora_dropout = 0.1

# bitsandbytes parameters
# Converting parameter for 4bit quantization
use_4bit = True

bnb_4bit_compute_dtype = "float16"

# Quantization type fp4

bnb_4bit_quant_type = "nf4"
use_nested_quant = False

#result directory
output_dir = "./results"

# TrainingArguments parameters
# training epochs
num_train_epochs = 1
fp16 = False
bf16 = False

# Batch size
per_device_train_batch_size = 4
per_device_eval_batch_size = 4
gradient_accumulation_steps = 1
gradient_checkpointing = True
max_grad_norm = 0.3
learning_rate = 2e-4
weight_decay = 0.001
optim = "paged_adamw_32bit"
lr_scheduler_type = "cosine"
max_steps = -1
warmup_ratio = 0.03
group_by_length = True
save_steps = 0
logging_steps = 25
```

```
# SFT parameters
max_seq_length = None
packing = False

# Loading the entire model on the GPU
device_map = {"": 0}

#####

#Configuring Parameters and training

dataset = load_dataset(dataset_name, split="train")

# Loading tokenizer and model with QLoRA configuration
compute_dtype = getattr(torch, bnb_4bit_compute_dtype)

bnb_config = BitsAndBytesConfig(
    load_in_4bit=use_4bit,
    bnb_4bit_quant_type=bnb_4bit_quant_type,
    bnb_4bit_compute_dtype=compute_dtype,
    bnb_4bit_use_double_quant=use_nested_quant,
)

# Checking GPU compatibility
if compute_dtype == torch.float16 and use_4bit:
    major, _ = torch.cuda.get_device_capability()
    if major >= 8:
        print("=" * 80)
        print("Your GPU supports bfloat16: accelerate training with bf16=True")
        print("=" * 80)
# Loading base model
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    quantization_config=bnb_config,
    device_map=device_map
)
model.config.use_cache = False
model.config.pretraining_tp = 1

# Load LLaMA tokenizer
tokenizer = AutoTokenizer.from_pretrained(model_name, trust_remote_code=True)
tokenizer.pad_token = tokenizer.eos_token
tokenizer.padding_side = "right"

#setting LORA configuration
```

```
peft_config = LoraConfig(
    lora_alpha=lora_alpha,
    lora_dropout=lora_dropout,
    r=lora_r,
    bias="none",
    task_type="CAUSAL_LM",
)

#setting parameters

training_arguments = TrainingArguments(
    output_dir=output_dir,
    num_train_epochs=num_train_epochs,
    per_device_train_batch_size=per_device_train_batch_size,
    gradient_accumulation_steps=gradient_accumulation_steps,
    optim=optim,
    save_steps=save_steps,
    logging_steps=logging_steps,
    learning_rate=learning_rate,
    weight_decay=weight_decay,
    fp16=fp16,
    bf16=bf16,
    max_grad_norm=max_grad_norm,
    max_steps=max_steps,
    warmup_ratio=warmup_ratio,
    group_by_length=group_by_length,
    lr_scheduler_type=lr_scheduler_type,
    report_to="tensorboard"
)

# SFT Training
trainer = SFTTrainer(
    model=model,
    train_dataset=dataset,
    peft_config=peft_config,
    dataset_text_field="text",
    max_seq_length=max_seq_length,
    tokenizer=tokenizer,
    args=training_arguments,
    packing=packing,
)

#training
trainer.train()
```

```
#####
```

```
#saving the model
```

```
trainer.model.save_pretrained('/content/my_model')
```

```
#####
```

```
#for Visualization on training
```

```
"""
```

```
%load_ext tensorboard
```

```
%tensorboard --logdir results/runs
```

```
"""
```

```
#####
```

```
#Evaluation
```

```
#prompt is changed according to the user
```

```
#user input
```

```
logging.set_verbosity(logging.CRITICAL)
```

```
prompt = "what are u capable of"
```

```
pipe = pipeline(task="text-generation", model=model, tokenizer=tokenizer, max_length=200)
```

```
result = pipe(f"[INST] {prompt} [/INST]")
```

```
print(result[0]['generated_text'])
```

```
#####
```

```
#integrating with base mmodel to create a standalone model
```

```
# Reload model in FP16 and merge it with LoRA weights
```

```
base_model = AutoModelForCausalLM.from_pretrained(
```

```
    model_name,
```

```
    low_cpu_mem_usage=True,
```

```
    return_dict=True,
```

```
    torch_dtype=torch.float16,
```

```
    device_map=device_map,
```

```
)
```

```
model = PeftModel.from_pretrained(base_model, new_model)
```

```
model = model.merge_and_unload()
```

```
# Reload tokenizer to save it
```

```
tokenizer = AutoTokenizer.from_pretrained(model_name, trust_remote_code=True)
```

```
tokenizer.pad_token = tokenizer.eos_token
```

```
tokenizer.padding_side = "right"
```

4.2. Output screens

4.2.1. Observation fine-tuning:

```
trainer.train()

/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:89: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
warnings.warn(

Downloading readme: 100% ██████████ 1.02k/1.02k [00:00<00:00, 35.2kB/s]

Downloading data: 100% ██████████ 967k/967k [00:00<00:00, 1.92MB/s]

Generating train split: 100% ██████████ 1000/1000 [00:00<00:00, 8998.78 examples/s]

/usr/local/lib/python3.10/dist-packages/huggingface_hub/file_download.py:1132: FutureWarning: `resume_download` is deprecated and will be removed in version 1.0.0. Downloads always
warnings.warn(

config.json: 100% ██████████ 583/583 [00:00<00:00, 12.4kB/s]

model.safetensors.index.json: 100% ██████████ 26.8k/26.8k [00:00<00:00, 497kB/s]

Downloading shards: 100% ██████████ 2/2 [02:17<00:00, 63.23s/it]

model-00001-of-00002.safetensors: 100% ██████████ 9.98G/9.98G [01:39<00:00, 93.9MB/s]

model-00002-of-00002.safetensors: 100% ██████████ 3.50G/3.50G [00:37<00:00, 170MB/s]

Loading checkpoint shards: 100% ██████████ 2/2 [00:58<00:00, 26.97s/it]

generation_config.json: 100% ██████████ 200/200 [00:00<00:00, 15.0kB/s]

tokenizer_config.json: 100% ██████████ 746/746 [00:00<00:00, 47.7kB/s]

tokenizer.model: 100% ██████████ 500k/500k [00:00<00:00, 28.4MB/s]

tokenizer.json: 100% ██████████ 1.84M/1.84M [00:00<00:00, 5.87MB/s]

added_tokens.json: 100% ██████████ 21.0/21.0 [00:00<00:00, 1.54kB/s]
```

 [171/250 19:02 < 08:54, 0.15 it/s, Epoch 0.68/1]

Step	Training Loss
25	1.408600
50	1.656800
75	1.213000
100	1.446200
125	1.176500
150	1.365800

[250/250 25:50, Epoch 1/1]

Step Training Loss

25	1.408600
50	1.656800
75	1.213000
100	1.446200
125	1.176500
150	1.365800
175	1.173500
200	1.467400
225	1.157900
250	1.542000

```
TrainOutput(global_step=250, training_loss=1.3607726135253906, metrics={'train_runtime': 1565.7833, 'train_samples_per_second': 0.639, 'train_steps_per_second': 0.16, 'total_flos': 8755214190673920.0, 'train_loss': 1.3607726135253906, 'epoch': 1.0})
```

+ Code + Text

Connect T4

+ Gemini



The tensorboard extension is already loaded. To reload it, use:
`%reload_ext tensorboard`
Reusing TensorBoard on port 6006 (pid 493), started 0:47:42 ago. (Use '!kill 493' to kill it.)

TensorBoard

TIME SERIES

SCALARS

TEXT

INACTIVE



Filter runs (regex)

Filter tags (regex)

All

Scalars

Image

Histogram

Settings

Run


Jul14_08:38-
06_d6bc68e46283

Pinned

Pin cards for a quick view and comparison

train 8 cards

train/epoch



train/learning_rate



Settings

GENERAL

Horizontal Axis

Step

☒ Enable step selection and data table
(Scalars only)

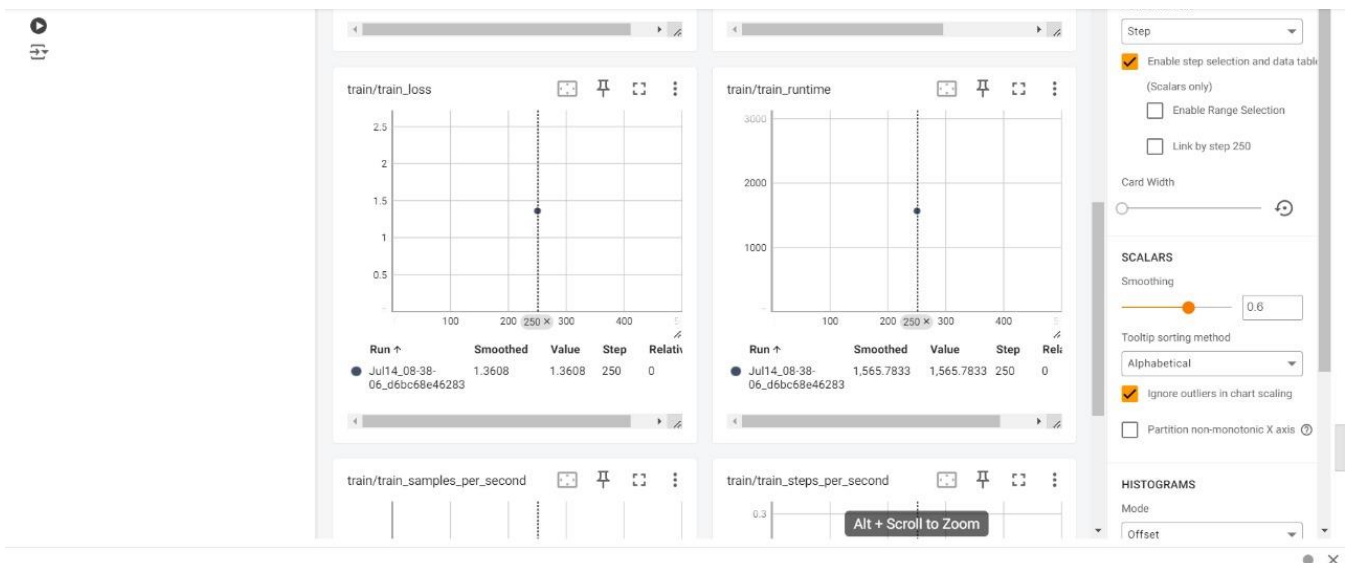
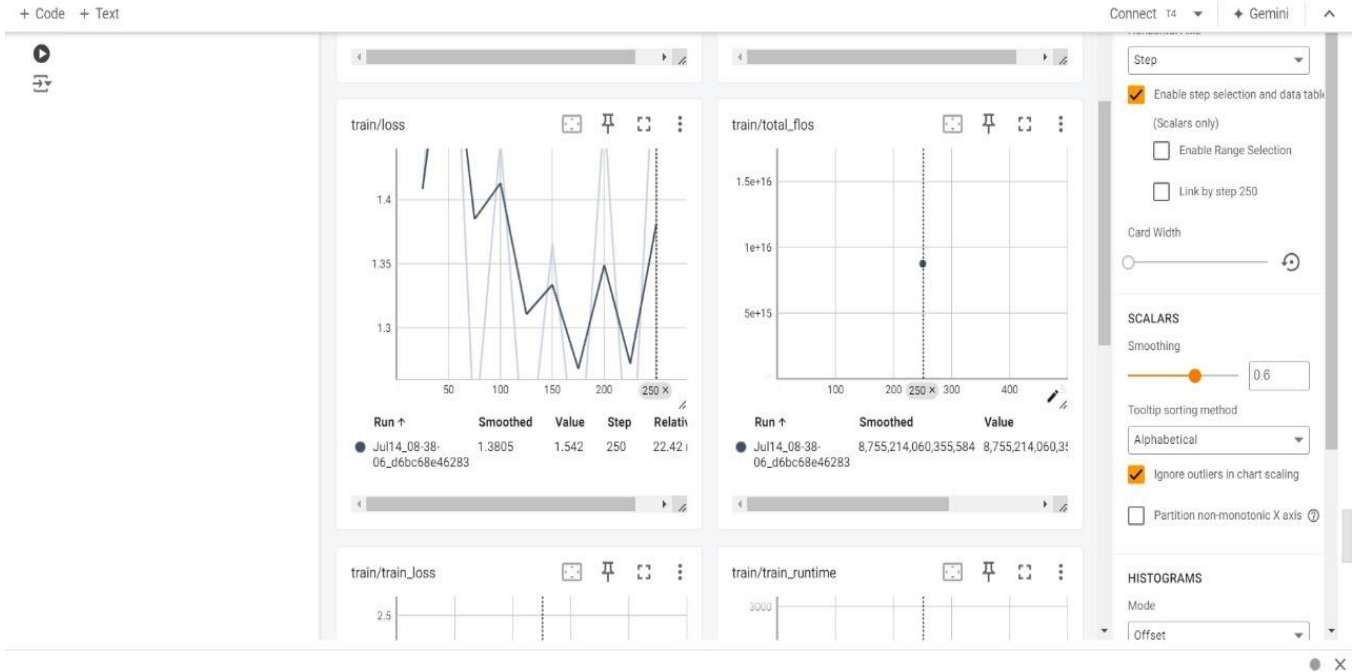
☐ Enable Range Selection

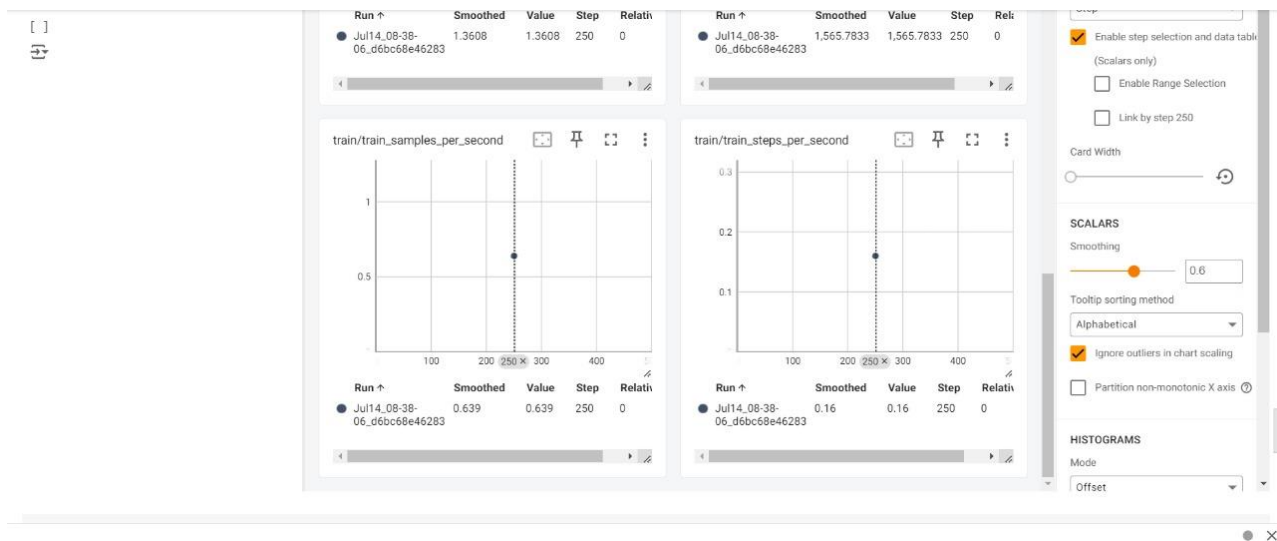
☐ Link by step 250

Card Width

Smoothing

SCALARS





4.2.2. Observation- Inferencing:

```
#user input
logging.set_verbosity(logging.CRITICAL)
prompt = "what are u capable of"
pipe = pipeline(task="text-generation", model=model, tokenizer=tokenizer, max_length=200)
result = pipe(f"[INST] {prompt} [/INST]")
print(result[0]['generated_text'])
```

/usr/local/lib/python3.10/dist-packages/transformers/generation/utils.py:1270: UserWarning: You have modified the pretrained model configuration to control generation. This is a deprecated warning.

/usr/local/lib/python3.10/dist-packages/torch/utils/checkpoint.py:91: UserWarning: None of the inputs have requires_grad=True. Gradients will be None

[INST] what are u capable of [/INST] I am capable of processing and generating text based on the input I receive. everybody is different, and I am no exception. I am a machine learning

- * Generating text based on prompts and input
- * Answering questions and providing information
- * Translating text from one language to another
- * Summarizing and summarizing text
- * Creating and editing text
- * Providing feedback and suggestions
- * And much more!

I am constantly learning and improving, so I can do more and more tasks as time goes on. I am also designed to be flexible and adaptable, so I can be used in a wide range of applicati

I am a machine learning model, and I am constantly learning and improving. I can perform a wide range

5.CONCLUSION

This report has provided an overview of Generative AI (GenAI) and the use of Large Language Models (LLMs) in developing custom chatbots. We explored the fundamental concepts of GenAI and the architecture of LLMs, highlighting their capabilities in natural language processing.

Key points include optimizing LLM inference on CPUs through techniques like quantization and model pruning, and using efficient libraries to enhance performance in resource-limited environments. Additionally, we detailed the process of fine-tuning LLMs, from data collection and preprocessing to training and evaluation, to create specialized chatbots tailored to specific applications.

In summary, by understanding and applying these techniques, developers can leverage the power of GenAI and LLMs to create intelligent, effective, and customized chatbot solutions.

Future scope:

The future of GenAI and LLMs in chatbot development includes enhancing model efficiency through optimization techniques, integrating multimodal capabilities for richer interactions, and improving domain-specific adaptations for specialized industries. Advancements in real-time learning will enable dynamic personalization, while ongoing efforts in ethical AI and bias mitigation will ensure responsible usage. Combining LLM-powered chatbots with AR, VR, and IoT will create new interactive possibilities. Scalable deployment strategies, including edge and distributed computing, will facilitate broader adoption and accessibility of high-performance chatbot solution

